

LEONID CHEKIN

PYTHON AI AND RAG DEVELOPER | BACKEND

Telegram: @leonchek97 | GitHub: github.com/leonidchekin55 | Portfolio: leonid-portfolio.onrender.com

Professional Summary

Python developer with commercial experience since 2022 across backend, integrations, and applied AI. Builds document ingestion and retrieval workflows for PDF, DOCX, and TXT, with chunking, embeddings, vector search, and cited passages. Uses Ollama for local generative RAG and FastAPI for service APIs.

Technical Skills

Languages: Python, SQL, JavaScript

Applied AI: LLM integrations, RAG, document chunking, embeddings, semantic search, Ollama

Vector and data: Qdrant, pgvector, PostgreSQL, SQLite

Backend: FastAPI, REST APIs, SQLAlchemy, Alembic, Pydantic, WebSockets

Delivery and integrations: Docker, Linux, Git, pytest, Celery, Redis, Telegram, GitHub Actions

Professional Experience

Python Backend Developer | Commercial development | 2022-present

- Develop Python backend services, REST APIs, data models, and third-party integrations.
- Handle validation, access control, background jobs, retries, error handling, and diagnostics using FastAPI, SQLAlchemy, PostgreSQL, Redis, Celery, and Docker.

Quantitative and Algorithmic Trading | Approximately 3.5 years

- Worked with historical market data, time-series analysis, quantitative calculations, and workflow automation.

Payment Integration | Tilda, T-Bank Acquiring, ATOL Online

- Implemented a customer-accepted flow for payments, status checks, basket parameters, and online receipts.

AI and Business Process Automation

- Built assistant workflows for incoming data, semantic deduplication, and CRM/Telegram integrations.

Selected Projects

Knowledge Assistant | RAG Demo. Ingests PDF/DOCX/TXT, chunks text, retrieves relevant passages, and returns source references with page numbers and chat history. Cloud demo uses extractive search; full RAG with Ollama and Qdrant runs locally. Stack: FastAPI, React, Qdrant, embeddings, PostgreSQL.

AI News RSS Bot. Collects feeds, removes semantic duplicates, and groups related stories for Telegram. Stack: Python, RSS, Ollama, embeddings, PostgreSQL/pgvector.

Pulseboard | SaaS Backend Demo. Multi-tenant API with JWT/RBAC, project and task workflows, signed idempotent webhooks, Celery jobs, WebSocket events, and an isolated public board. Stack: FastAPI, SQLAlchemy, PostgreSQL, Redis, Docker.

AI Engineer Portfolio | Learning project. Six FastAPI demos for structured intake, document processing, cited RAG, evals, and approval-gated actions. Public mode is mock-only. Stack: PostgreSQL/pgvector, Redis/Celery, FastEmbed, Docker.